



```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # Importing dataset
df = pd.read_csv('/content/Telco_Customer_Churn.csv')
df
```

```
Out[2]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneS
0	7590-VHVEG	Female	0	Yes	No	1	
1	5575-GNVDE	Male	0	No	No	34	
2	3668-QPYBK	Male	0	No	No	2	
3	7795-CFOCW	Male	0	No	No	45	
4	9237-HQITU	Female	0	No	No	2	
...
7038	6840-RESVB	Male	0	Yes	Yes	24	
7039	2234-XADUH	Female	0	Yes	Yes	72	
7040	4801-JZAZL	Female	0	Yes	Yes	11	
7041	8361-LTMKD	Male	1	Yes	No	4	
7042	3186-AJIEK	Male	0	No	No	66	

7043 rows × 21 columns

EDA and Data exploration

```
In [3]: df.shape
```

```
Out[3]: (7043, 21)
```

```
In [4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   customerID        7043 non-null    object  
 1   gender             7043 non-null    object  
 2   SeniorCitizen     7043 non-null    int64  
 3   Partner            7043 non-null    object  
 4   Dependents         7043 non-null    object  
 5   tenure             7043 non-null    int64  
 6   PhoneService       7043 non-null    object  
 7   MultipleLines      7043 non-null    object  
 8   InternetService    7043 non-null    object  
 9   OnlineSecurity     7043 non-null    object  
 10  OnlineBackup       7043 non-null    object  
 11  DeviceProtection  7043 non-null    object  
 12  TechSupport        7043 non-null    object  
 13  StreamingTV        7043 non-null    object  
 14  StreamingMovies    7043 non-null    object  
 15  Contract           7043 non-null    object  
 16  PaperlessBilling   7043 non-null    object  
 17  PaymentMethod      7043 non-null    object  
 18  MonthlyCharges    7043 non-null    float64 
 19  TotalCharges       7043 non-null    object  
 20  Churn              7043 non-null    object  
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

- 7,043 rows and 21 columns form a medium-sized churn classification dataset.
- Most features are categorical, making encoding critical.
- The target `Churn` is object it has to converted into binary variable.
- `TotalCharges` has an incorrect data type and needs conversion.
- `SeniorCitizen` is categorical despite numeric encoding.

```
In [5]: # Exploring target variable
print('Values in Churn column', df['Churn'].unique())
print('Value counts:', df['Churn'].value_counts())
```

```
Values in Churn column ['No' 'Yes']
Value counts: Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

- The target variable `Churn` is binary with values `Yes` and `No`.
- Non-churned customers dominate the dataset (5,174 vs 1,869), indicating class imbalance.

- Roughly 26% of customers have churned, making accuracy an unreliable evaluation metric.
- Recall and F1-score for the churn class will be more informative than overall accuracy.

```
In [6]: for col in df.columns:
    print(f'{col}:{df[col].unique()}')


customerID: ['7590-VHVEG' '5575-GNVDE' '3668-QPYBK' ... '4801-JZAZL' '8361-LTMK
D'
'3186-AJIEK']
gender: ['Female' 'Male']
SeniorCitizen: [0 1]
Partner: ['Yes' 'No']
Dependents: ['No' 'Yes']
tenure: [ 1 34  2 45  8 22 10 28 62 13 16 58 49 25 69 52 71 21 12 30 47 72 17 27
5 46 11 70 63 43 15 60 18 66  9  3 31 50 64 56  7 42 35 48 29 65 38 68
32 55 37 36 41  6  4 33 67 23 57 61 14 20 53 40 59 24 44 19 54 51 26  0
39]
PhoneService: ['No' 'Yes']
MultipleLines: ['No phone service' 'No' 'Yes']
InternetService: ['DSL' 'Fiber optic' 'No']
OnlineSecurity: ['No' 'Yes' 'No internet service']
OnlineBackup: ['Yes' 'No' 'No internet service']
DeviceProtection: ['No' 'Yes' 'No internet service']
TechSupport: ['No' 'Yes' 'No internet service']
StreamingTV: ['No' 'Yes' 'No internet service']
StreamingMovies: ['No' 'Yes' 'No internet service']
Contract: ['Month-to-month' 'One year' 'Two year']
PaperlessBilling: ['Yes' 'No']
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
'Credit card (automatic)']
MonthlyCharges: [29.85 56.95 53.85 ... 63.1 44.2 78.7 ]
TotalCharges: ['29.85' '1889.5' '108.15' ... '346.45' '306.6' '6844.5']
Churn: ['No' 'Yes']
```

```
In [7]: # Fixing data correctness

# Dropping 'customerID' as it adds zero predictive value to the model
df.drop(columns=['customerID'], inplace=True)
print('Verifying column drop:', ('customerID' in df.columns))

# Converting 'TotalCharges' to number as it should be
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
print(`TotalCharges` to number:, df['TotalCharges'].dtypes)

# Customers with zero tenure has to have no billing history, so TotalCharges is
df.loc[df['tenure']==0, 'TotalCharges'] = 0
print('Cheking TotalCharges is set to 0:', (df.loc[df['tenure'] == 0, 'TotalCh

# Converting Churn to binary
df['Churn'] = df['Churn'].map({'Yes': 1, 'No': 0})
```

```
print('Churn to binary:', df['Churn'].unique(), 'Type:', df['Churn'].dtypes)
```

```
Verifying column drop: False
`TotalCharges` to number: float64
Cheking TotalCharges is set to 0: True
Churn to binary: [0 1] Type: int64
```

- Non-informative identifier `customerID` was removed to prevent noise and leakage.
- `TotalCharges` was converted to a numeric feature and corrected for zero-tenure customers using business logic.
- The target variable `Churn` was mapped to a binary format suitable for classification models.

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   gender            7043 non-null    object 
 1   SeniorCitizen     7043 non-null    int64  
 2   Partner           7043 non-null    object 
 3   Dependents        7043 non-null    object 
 4   tenure            7043 non-null    int64  
 5   PhoneService      7043 non-null    object 
 6   MultipleLines     7043 non-null    object 
 7   InternetService   7043 non-null    object 
 8   OnlineSecurity    7043 non-null    object 
 9   OnlineBackup       7043 non-null    object 
 10  DeviceProtection  7043 non-null    object 
 11  TechSupport       7043 non-null    object 
 12  StreamingTV       7043 non-null    object 
 13  StreamingMovies   7043 non-null    object 
 14  Contract          7043 non-null    object 
 15  PaperlessBilling  7043 non-null    object 
 16  PaymentMethod     7043 non-null    object 
 17  MonthlyCharges   7043 non-null    float64
 18  TotalCharges      7043 non-null    float64
 19  Churn             7043 non-null    int64  
dtypes: float64(2), int64(3), object(15)
memory usage: 1.1+ MB
```

```
In [9]: df.isnull().sum()
```


Out[9]:

	0
gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	0
Churn	0

dtype: int64

In [10]:

```
# Target distribution
sns.countplot(x='Churn', data=df)
plt.title('Churn Distribution')
plt.show()
```




- The target variable is moderately imbalanced, with non-churn customers significantly outnumbering churners, making accuracy an unreliable evaluation metric.
- Approximately one-quarter of customers churn, indicating the need to prioritize recall and F1-score for the churn class during modeling.

```
In [11]: # Target vs Categorical features (Churn rate)
cat_cols = df.select_dtypes(include='object').columns

fig, axes = plt.subplots(nrows=5, ncols=3, figsize=(18, 20))
axes = axes.flatten()

for ax, col in zip(axes, cat_cols):
    churn_rate = df.groupby(col)['Churn'].mean().sort_values(ascending=False)
    churn_rate.plot(kind='bar', ax=ax)
    ax.set_title(f'Churn Rate by {col}')
    ax.set_ylabel('Churn Rate')
    ax.set_xlabel(col)
    ax.tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()
```



- Gender shows nearly identical churn rates, indicating it has minimal predictive value.
- Customers without partners or dependents churn significantly more, suggesting household stability reduces churn risk.
- PhoneService and MultipleLines show weak churn separation, implying limited standalone predictive power.
- Fiber optic internet users churn at much higher rates than DSL or no-internet customers, making InternetService a strong churn driver.


- Lack of OnlineSecurity, OnlineBackup, DeviceProtection, or TechSupport is consistently associated with very high churn rates.
- Customers with no internet service exhibit the lowest churn, indicating bundled internet services drive churn risk.
- Streaming services (TV and Movies) show moderate churn differences, suggesting secondary influence rather than primary drivers.
- Month-to-month contracts have drastically higher churn than one- or two-year contracts, making Contract the strongest categorical predictor.
- Paperless billing customers churn more, indicating a behavioral or billing-related churn signal.
- Electronic check users churn far more than other payment methods, highlighting PaymentMethod as a high-risk segment.

```
In [12]: # Target vs Numerical features (Distribution comparison)
num_cols = ['tenure', 'MonthlyCharges', 'TotalCharges']

fig, axes = plt.subplots(1, 3, figsize=(18, 5))

for ax, col in zip(axes, num_cols):
    sns.boxplot(x='Churn', y=col, data=df, ax=ax)
    ax.set_title(f'{col} vs Churn')

plt.tight_layout()
plt.savefig('numeric_features_boxplots.png', dpi=300, bbox_inches='tight')
plt.show()
```




- Churned customers have significantly lower tenure, confirming tenure as a strong inverse predictor of churn.
- Customers who churn tend to have higher monthly charges, indicating pricing pressure contributes to churn risk.
- TotalCharges is substantially lower for churned customers, reflecting shorter customer lifetimes rather than lower spending behavior.
- Heavy overlap in MonthlyCharges distributions suggests churn is not

driven by price alone but by interactions with contract and service features.

- The strong tenure–TotalCharges relationship implies redundancy and requires cautious interpretation of feature importance.

```
In [13]: # Tenure bucketing
df_tenure_bucket = df.copy()
df_tenure_bucket['tenure_bucket'] = pd.cut(
    df['tenure'],
    bins=[0, 12, 24, 36, 48, 60, 72],
    right=False
)

plt.figure(figsize=(8,4))
sns.barplot(data=df_tenure_bucket, x='tenure_bucket', y='Churn', hue='tenure_b
plt.title('Churn Rate by Tenure Bucket')
plt.ylabel('Chrun Rate')
plt.xlabel('Tenure Buckets')
plt.show()
```




- Churn rate is highest in the first year of tenure and declines monotonically as customer tenure increases.
- Nearly half of customers churn within the first 12 months, indicating early tenure as the highest-risk period.
- Long-tenure customers (60+ months) exhibit very low churn, reflecting strong customer loyalty.
- The clear downward trend confirms tenure as the single strongest

predictor of churn.

- Non-linear churn decay across tenure buckets justifies the use of tree-based models over linear approaches.

```
In [14]: # Correlation and Redundancy
corr = df[['tenure', 'MonthlyCharges', 'TotalCharges', 'Churn']].corr()

sns.heatmap(corr, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.xticks(rotation=45)
plt.savefig('Correlation Matrix.png', dpi=300, bbox_inches='tight')
plt.show()
```



- TotalCharges is strongly correlated with tenure, confirming redundancy driven by customer lifetime rather than independent spending behavior.
- The weak linear correlations with churn suggest that non-linear interactions, rather than single-feature effects, dominate churn

behavior.

- High correlation between tenure and TotalCharges implies feature importance from tree models must be interpreted cautiously.

Random Forest Classification Model

```
In [15]: !pip install category_encoders
```

```
Collecting category_encoders
  Downloading category_encoders-2.9.0-py3-none-any.whl.metadata (7.9 kB)
Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.12/dist-
packages (from category_encoders) (2.0.2)
Requirement already satisfied: pandas>=1.0.5 in /usr/local/lib/python3.12/dist-
packages (from category_encoders) (2.2.2)
Requirement already satisfied: patsy>=0.5.1 in /usr/local/lib/python3.12/dist-p
ackages (from category_encoders) (1.0.2)
Requirement already satisfied: scikit-learn>=1.6.0 in /usr/local/lib/python3.1
2/dist-packages (from category_encoders) (1.6.1)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.12/dist-p
ackages (from category_encoders) (1.16.3)
Requirement already satisfied: statsmodels>=0.9.0 in /usr/local/lib/python3.12/
dist-packages (from category_encoders) (0.14.6)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python
3.12/dist-packages (from pandas>=1.0.5->category_encoders) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-p
ackages (from pandas>=1.0.5->category_encoders) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dis
t-packages (from pandas>=1.0.5->category_encoders) (2025.3)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-
packages (from scikit-learn>=1.6.0->category_encoders) (1.5.3)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.1
2/dist-packages (from scikit-learn>=1.6.0->category_encoders) (3.6.0)
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.12/dis
t-packages (from statsmodels>=0.9.0->category_encoders) (25.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-pac
kes (from python-dateutil>=2.8.2->pandas>=1.0.5->category_encoders) (1.17.0)
  Downloading category_encoders-2.9.0-py3-none-any.whl (85 kB)
    0.0/85.9 kB ? eta -----+
    85.9/85.9 kB 3.0 MB/s eta 0:00:00
```

```
Installing collected packages: category_encoders
Successfully installed category_encoders-2.9.0
```

```
In [16]: # Importing sikit-learn libraries
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestClassifier
from category_encoders.target_encoder import TargetEncoder
from sklearn.metrics import classification_report, roc_auc_score, accuracy_scc
```

```
In [17]: # Spiting DataFrame into X and y (Features and Target)
X = df.drop(columns='Churn')
y = df['Churn']

# Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42, str
```

```
In [18]: # Identifying categorical columns and number columns
cat_cols = X_train.select_dtypes(include='object').columns
num_cols = X_train.select_dtypes(exclude='object').columns

# Adding Tarrget Encoding to categorical columns only
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', TargetEncoder(smoothing=10), cat_cols),
        ('num', 'passthrough', num_cols)
    ]
)
```

```
In [19]: # Random Forest with OOB
clf = RandomForestClassifier(n_estimators=200, class_weight='balanced', oob_sc
```

```
In [20]: # Pipeline
pipeline = Pipeline(
    steps=[
        ('preprocessor', preprocessor),
        ('model', clf)
    ]
)
```

```
In [21]: pipeline.fit(X_train, y_train)

# Evaluation
y_pred = pipeline.predict(X_test)
y_proba = pipeline.predict_proba(X_test)[:,1]


# Scores
oob_score = pipeline.named_steps['model'].oob_score_
print('OOB socre:', oob_score)
print('Accuracy Score:', accuracy_score(y_test, y_pred))
print('ROC-AUC:', roc_auc_score(y_test, y_proba))
print(classification_report(y_test, y_pred))
```

```
00B score: 0.7900248491302805
Accuracy Score: 0.7849538679914834
ROC-AUC: 0.8241210054509287
      precision    recall   f1-score   support
0           0.83     0.89     0.86     1035
1           0.62     0.48     0.54      374
accuracy          0.78     0.78     0.78     1409
macro avg       0.72     0.69     0.70     1409
weighted avg     0.77     0.78     0.78     1409
```

```
In [22]: # Computing confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Ploting heatmap
sns.heatmap(
    cm,
    annot=True,
    fmt='d',
    cmap='Blues',
    xticklabels=['No Churn', 'Churn'],
    yticklabels=['No Churn', 'Churn']
)

plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.tight_layout()
plt.show()
```



```
In [23]: thresholds = np.arange(0.2, 0.6, 0.05)

for t in thresholds:
    y_pred_t = (y_proba >= t).astype(int)
    print(f"\nThreshold: {t:.2f}")
    print(classification_report(y_test, y_pred_t))
```

Threshold: 0.20

	precision	recall	f1-score	support
0	0.92	0.67	0.77	1035
1	0.48	0.84	0.61	374
accuracy			0.71	1409
macro avg	0.70	0.75	0.69	1409
weighted avg	0.80	0.71	0.73	1409

Threshold: 0.25

	precision	recall	f1-score	support
0	0.90	0.71	0.79	1035
1	0.49	0.78	0.60	374
accuracy			0.73	1409
macro avg	0.69	0.74	0.70	1409
weighted avg	0.79	0.73	0.74	1409

Threshold: 0.30

	precision	recall	f1-score	support
0	0.89	0.77	0.82	1035
1	0.53	0.72	0.61	374
accuracy			0.76	1409
macro avg	0.71	0.75	0.72	1409
weighted avg	0.79	0.76	0.77	1409

Threshold: 0.35

	precision	recall	f1-score	support
0	0.87	0.81	0.84	1035
1	0.56	0.67	0.61	374
accuracy			0.77	1409
macro avg	0.72	0.74	0.73	1409
weighted avg	0.79	0.77	0.78	1409

Threshold: 0.40

	precision	recall	f1-score	support
0	0.86	0.84	0.85	1035
1	0.58	0.62	0.60	374
accuracy			0.78	1409
macro avg	0.72	0.73	0.73	1409
weighted avg	0.79	0.78	0.78	1409

Threshold: 0.45

	precision	recall	f1-score	support
0	0.84	0.87	0.85	1035
1	0.60	0.55	0.57	374

	accuracy		0.78	1409
macro avg	0.72	0.71	0.71	1409
weighted avg	0.78	0.78	0.78	1409

Threshold: 0.50

	precision	recall	f1-score	support
0	0.83	0.89	0.86	1035
1	0.62	0.49	0.54	374

	accuracy		0.78	1409
macro avg	0.72	0.69	0.70	1409
weighted avg	0.77	0.78	0.78	1409

Threshold: 0.55

	precision	recall	f1-score	support
0	0.82	0.91	0.86	1035
1	0.65	0.43	0.52	374

	accuracy		0.79	1409
macro avg	0.73	0.67	0.69	1409
weighted avg	0.77	0.79	0.77	1409

- A decision threshold of 0.35 was selected to balance churn recall and precision, increasing churn recall from 0.48 to 0.67 while maintaining stable overall performance.

```
In [24]: leaf_sizes = [5, 10, 25, 50]

for leaf in leaf_sizes:
    pipeline.named_steps['model'].set_params(min_samples_leaf=leaf)
    pipeline.fit(X_train, y_train)

    oob = pipeline.named_steps['model'].oob_score_
    y_proba = pipeline.predict_proba(X_test)[:, 1]

    print(f"min_samples_leaf={leaf} | OOB={oob:.3f} | ROC-AUC={roc_auc_score(y
```

```
min_samples_leaf=5 | OOB=0.778 | ROC-AUC=0.840
min_samples_leaf=10 | OOB=0.771 | ROC-AUC=0.844
min_samples_leaf=25 | OOB=0.760 | ROC-AUC=0.845
min_samples_leaf=50 | OOB=0.753 | ROC-AUC=0.844
```

- Increasing min_samples_leaf reduced model variance without affecting ranking performance, and a value of 25 was selected to maximize ROC-AUC while improving generalization.

```
In [25]: depths = [None, 10, 15, 20]

for d in depths:
    pipeline.named_steps['model'].set_params(max_depth=d)
    pipeline.fit(X_train, y_train)

    oob = pipeline.named_steps['model'].oob_score_
    y_proba = pipeline.predict_proba(X_test)[:, 1]

    print(f"max_depth={d} | OOB={oob:.3f} | ROC-AUC={roc_auc_score(y_test, y_p
max_depth=None | OOB=0.753 | ROC-AUC=0.844
max_depth=10 | OOB=0.753 | ROC-AUC=0.845
max_depth=15 | OOB=0.753 | ROC-AUC=0.844
max_depth=20 | OOB=0.753 | ROC-AUC=0.844
```

- Random Forest performance was insensitive to max_depth once regularized by min_samples_leaf, and a depth of 15 was selected as a safe upper bound without sacrificing ROC-AUC.

```
In [26]: # final model parameters
pipeline.named_steps['model'].set_params(
    min_samples_leaf=25,
    max_depth=15
)

pipeline.fit(X_train, y_train)

# Scores
oob_score = pipeline.named_steps['model'].oob_score_
print('OOB score:', oob_score)
print('Accuracy Score:', accuracy_score(y_test, y_pred))
print('ROC-AUC:', roc_auc_score(y_test, y_proba))
print(classification_report(y_test, y_pred))
```

```

OOB score: 0.7598509052183173
Accuracy Score: 0.7849538679914834
ROC-AUC: 0.8444354026195457
      precision    recall   f1-score   support
0           0.83     0.89     0.86     1035
1           0.62     0.48     0.54      374

accuracy          0.78     1409
macro avg       0.72     0.69     0.70     1409
weighted avg    0.77     0.78     0.78     1409

```

```


In [27]: # Apply chosen threshold
best_threshold = 0.35
y_pred_final = (y_proba >= best_threshold).astype(int)

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred_final)

# Plotting heatmap
sns.heatmap(
    cm,
    annot=True,
    fmt='d',
    cmap='Blues',
    xticklabels=['No Churn', 'Churn'],
    yticklabels=['No Churn', 'Churn']
)

plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.tight_layout()
plt.show()

```




```
In [28]: # Extracting trained Random Forest
rf_model = pipeline.named_steps['model']

# Getting feature names after preprocessing
cat_cols = pipeline.named_steps['preprocessor'].transformers_[0][2]
num_cols = pipeline.named_steps['preprocessor'].transformers_[1][2]
feature_names = list(cat_cols) + list(num_cols)

# Creating feature importance DataFrame
feature_importance = pd.DataFrame({
    'Feature': feature_names,
    'Importance': rf_model.feature_importances_
}).sort_values(by='Importance', ascending=False)

# Selecting top 10 features
top_10 = feature_importance.head(10)

# Plot
plt.figure(figsize=(8, 5))
sns.barplot(data=top_10, x='Importance', y='Feature', hue='Feature', palette='husl')
plt.gca().invert_yaxis()
plt.xlabel('Feature Importance')
plt.title('Top 10 Feature Importances (Random Forest)')
plt.tight_layout()
plt.show()
```



Final Insights & Project Conclusion

This project focused on predicting customer churn using the Telco Customer Churn dataset, which contains 7,043 customer records and a mix of categorical and numerical features after removing non-informative identifiers. Given the moderate class imbalance (approximately 26% churners), the problem was framed as a risk identification task rather than a pure accuracy optimization exercise.

Modeling Approach and Performance

A Random Forest classifier was selected due to its ability to capture complex, non-linear relationships and interactions between features. Categorical variables were handled using target encoding, which preserved ordinal churn risk information without exploding dimensionality. The model was trained using class weighting and evaluated with Out-of-Bag (OOB) scoring to estimate generalization performance without excessive cross-validation.

The final model achieved:

- OOB Score: ~0.76
- ROC-AUC: ~0.84

- Accuracy: ~0.78

Hyperparameter tuning showed that increasing `min_samples_leaf` reduced variance without materially affecting ranking performance, and the model was largely insensitive to `max_depth` once regularized, indicating stable generalization behavior.

Threshold Optimization and Business Interpretation

Rather than relying on the default 0.5 decision threshold, multiple thresholds were evaluated to balance false positives and false negatives. A threshold of 0.35 was selected, increasing churn recall from approximately 0.48 to 0.67 while maintaining acceptable precision. This choice reflects a business-driven decision: in churn prevention, missing a true churner (Type II error) is typically more costly than incorrectly flagging a loyal customer. Keeping the threshold external to the model ensures flexibility and allows the decision policy to adapt to changing business costs without retraining the model.

Final Conclusion

Overall, the analysis confirms that customer churn in this dataset is driven by early tenure risk, contract flexibility, internet service characteristics, payment behavior, and value-added services, rather than static demographic traits. The presence of strong non-linear patterns and feature interactions justified the use of a Random Forest model with target encoding and careful threshold optimization. The final solution demonstrates not only solid predictive performance but also clear alignment with real-world business decision-making, making it suitable for practical churn risk monitoring and targeted retention strategies.

In [28]: